

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, roc_auc_score
from sklearn.decomposition import PCA

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
```

```
In [2]: df=pd.read_csv('C:/Users/mmi/Downloads/titanic/train.csv')

df.head()
```

```
Out[2]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
In [3]: df.fillna(df.mean(),inplace=True)
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: PassengerId      0
Survived      0
Pclass      0
Name      0
Sex      0
Age      0
SibSp      0
Parch      0
Ticket      0
Fare      0
Cabin      687
Embarked      2
dtype: int64
```

```
In [5]: df.drop(['PassengerId','Name','Ticket','Cabin'], axis=1, inplace=True)
```

```
In [6]: df.head()
```

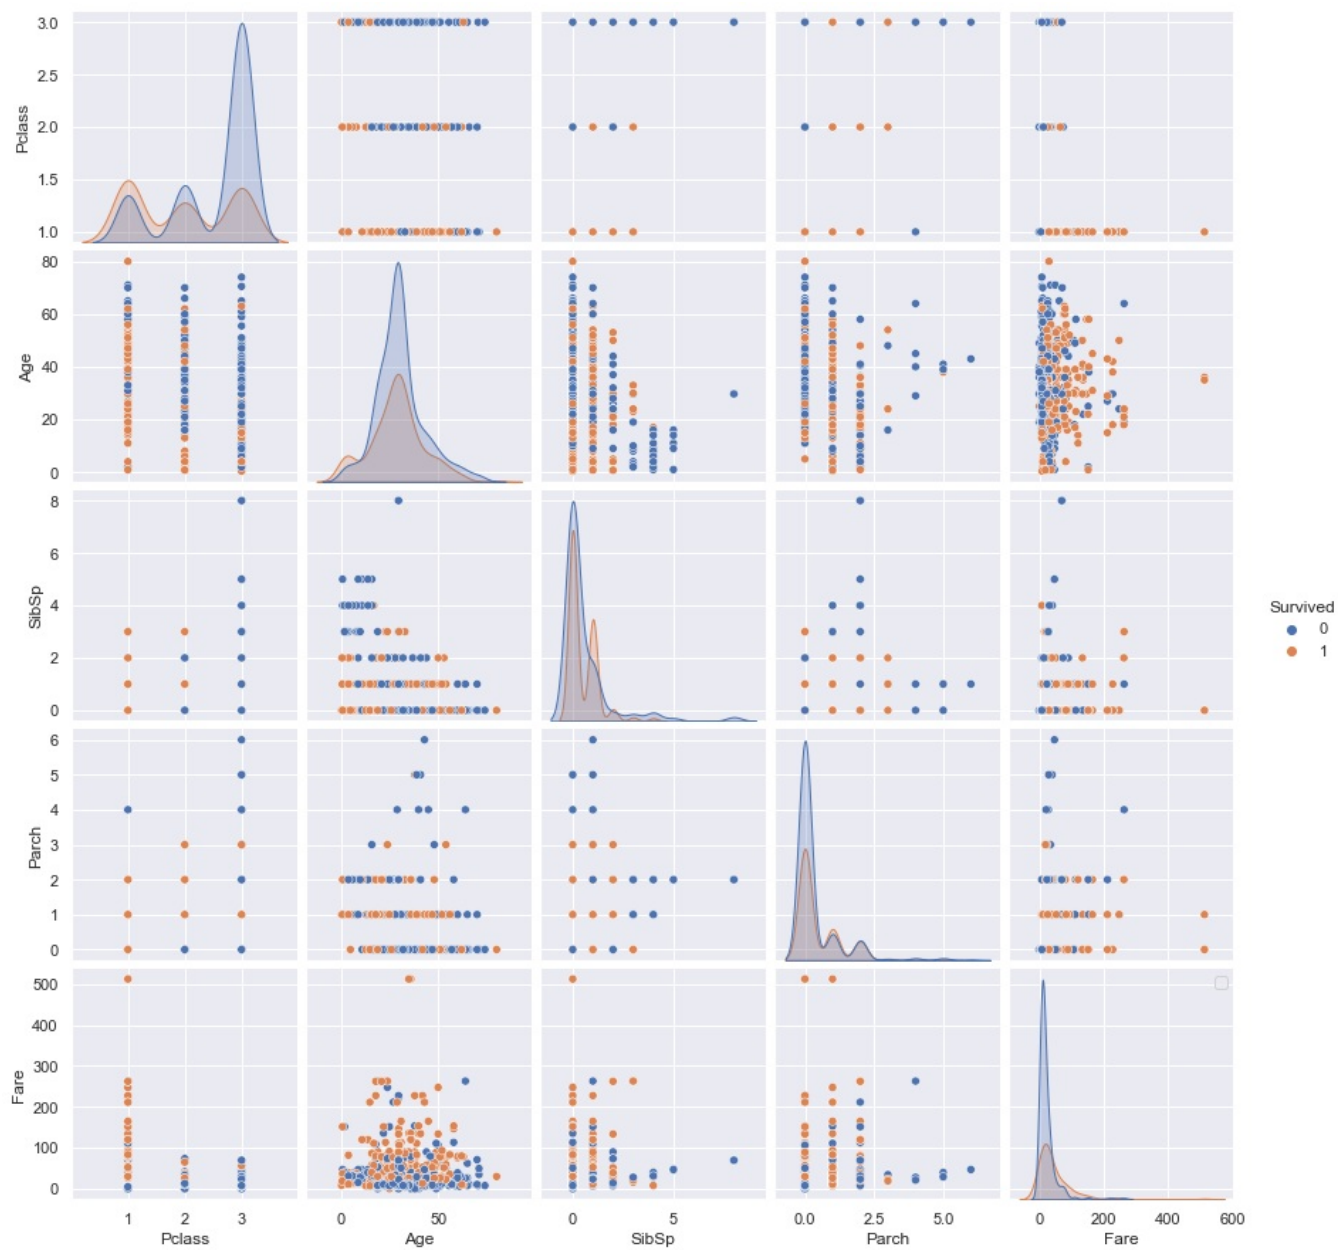
```
Out[6]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S


```
In [7]: sns.set()

cols_to_pairplot = df.columns[:11]
sns.pairplot(df[cols_to_pairplot], hue="Survived")
plt.legend()
plt.show()
```


No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.





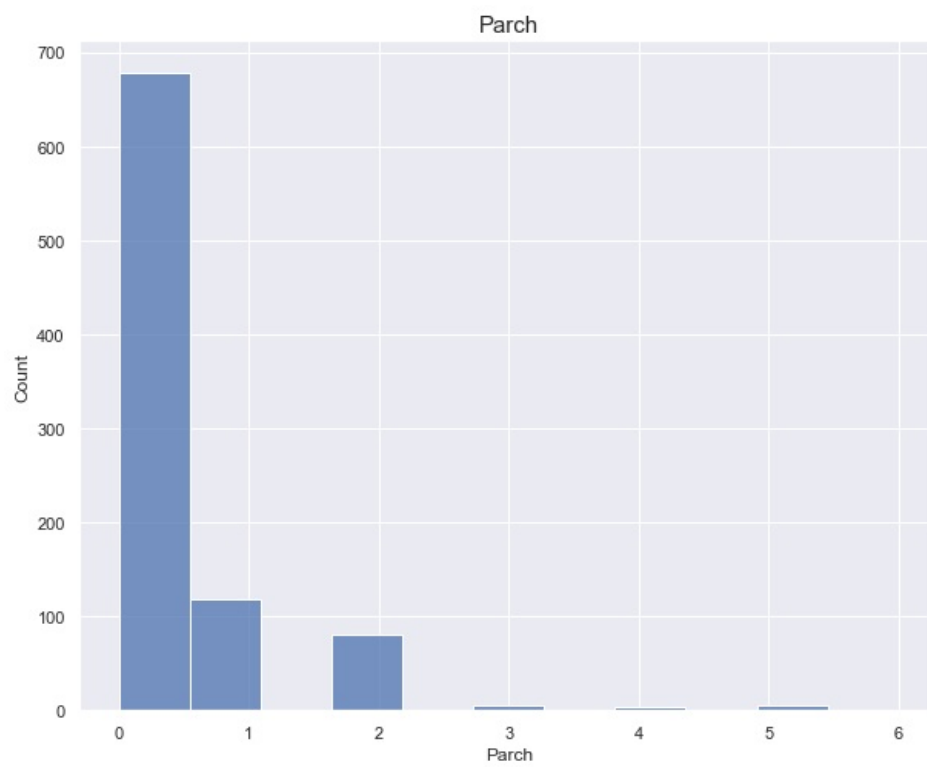
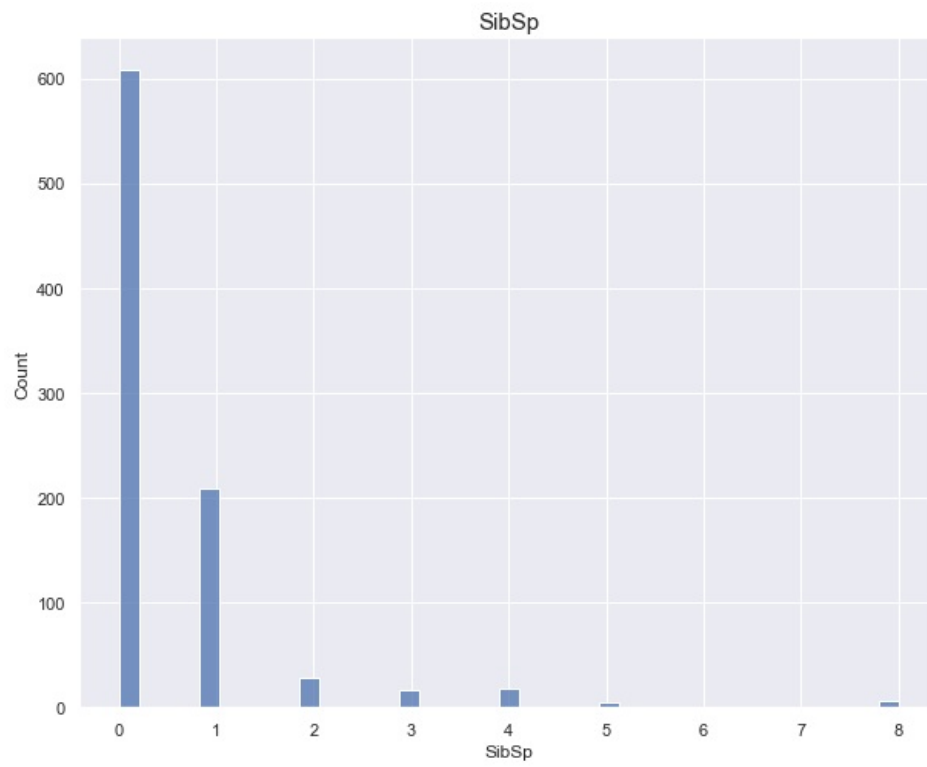
```
In [8]: plt.figure(figsize=(10,8))
sns.countplot(df["Survived"])
plt.title("Survived", size=15)
plt.show()
```





```
In [9]: for col in df.drop("Survived", axis=1).columns:  
        plt.figure(figsize=(10,8))  
        sns.histplot(df[col])  
        plt.title(f"{col}", size=15)  
        plt.show()
```

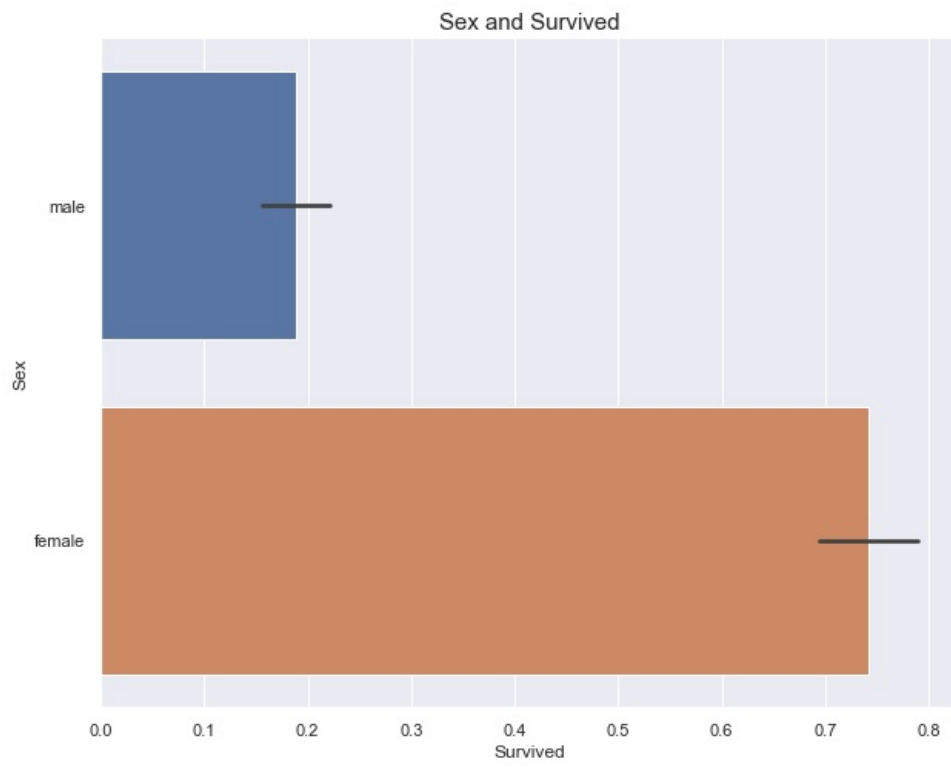
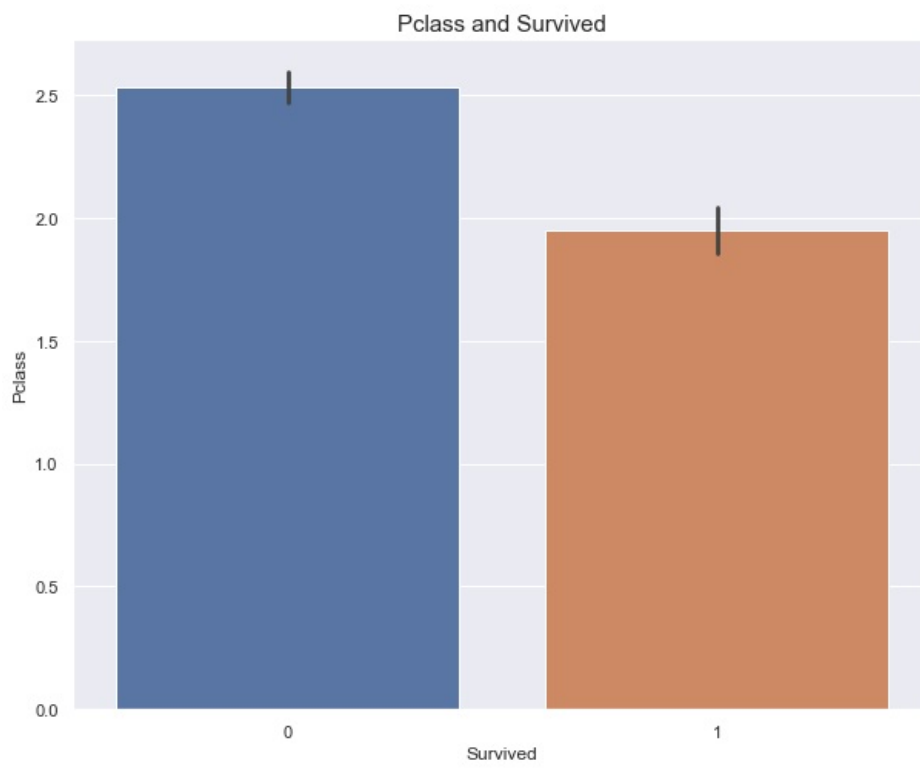


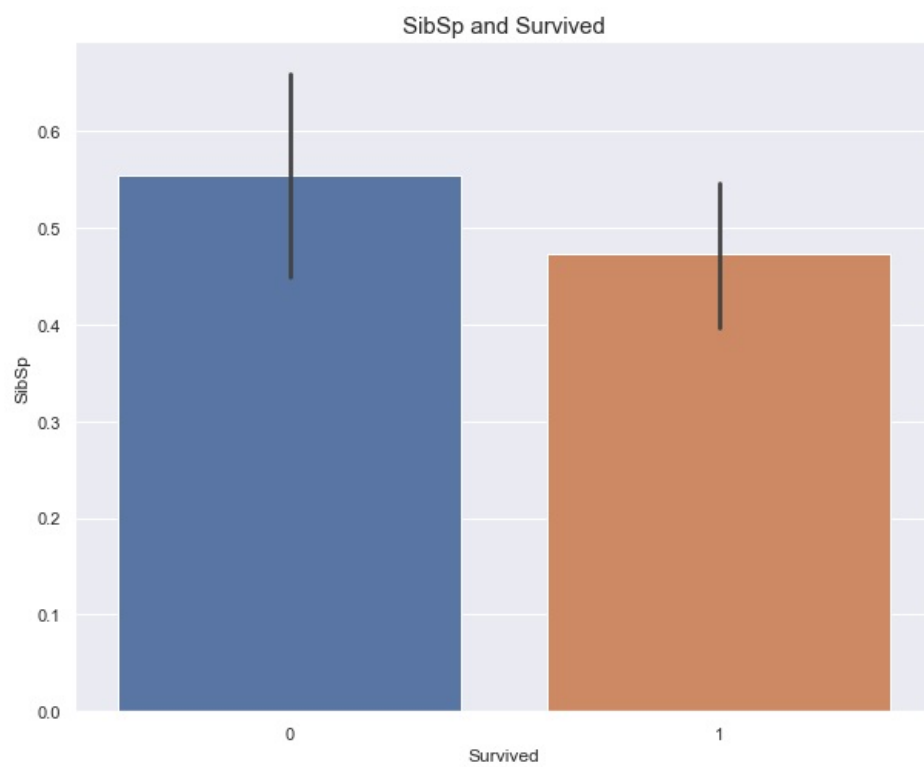




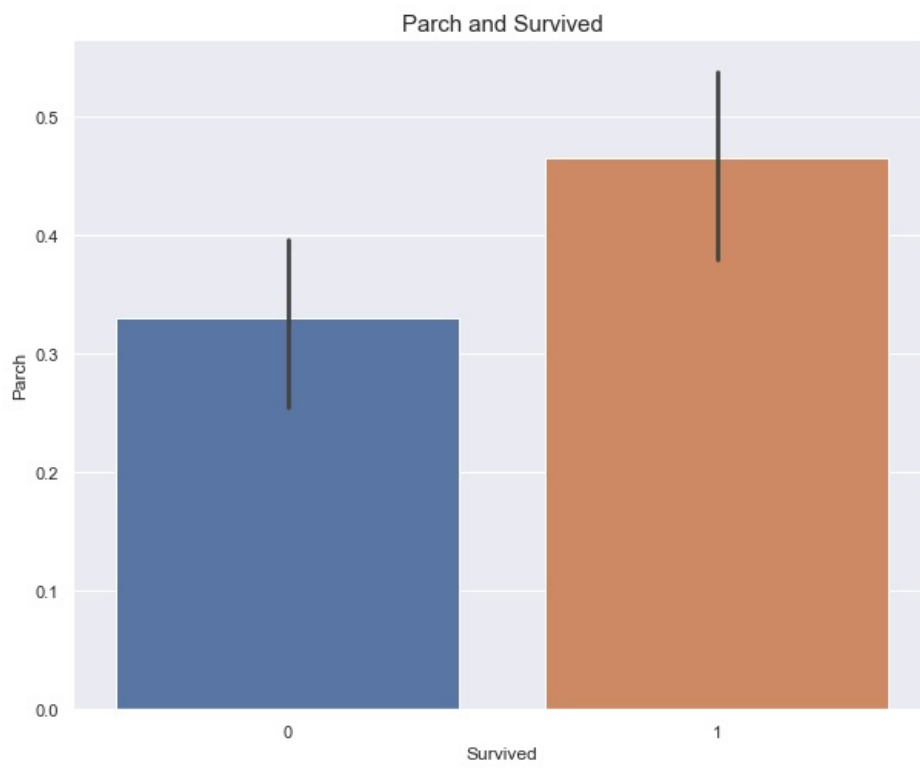


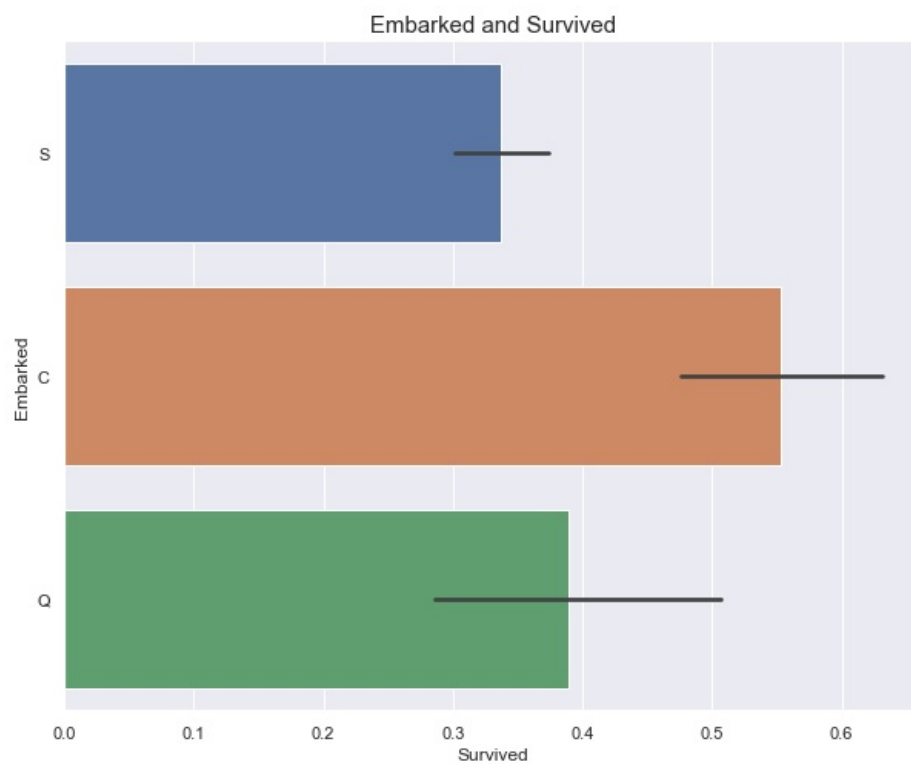


```
In [10]: for col in df.drop("Survived", axis=1).columns:
plt.figure(figsize=(10,8))
sns.barplot(x=df["Survived"], y=df[col])
plt.title(f"{col} and Survived", size=15)
plt.show()
```








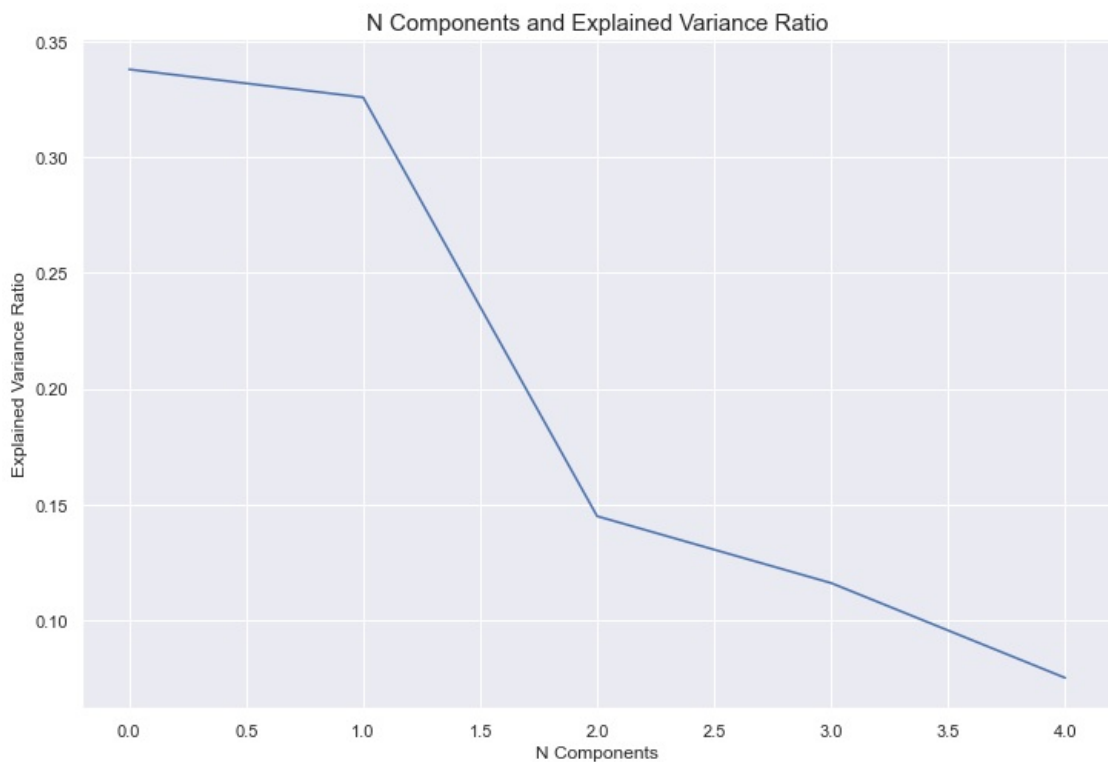


```
In [11]: plt.figure(figsize=(14,10))
sns.heatmap(df.corr(), cmap="Blues")
plt.title("Correlations Between Variables", size=16)
plt.show()
```



```
In [18]: X = df.drop("Survived", axis=1)
y = df["Survived"].replace({"male": 0, "female": 1})
X=df[['Pclass','Age','SibSp','Parch','Fare']]
```

```
In [19]: scaler = StandardScaler()
X = scaler.fit_transform(X)
pca = PCA()
pca.fit(X)
plt.figure(figsize=(12,8))
plt.plot(pca.explained_variance_ratio_)
plt.title("N Components and Explained Variance Ratio", size=15)
plt.xlabel("N Components")
plt.ylabel("Explained Variance Ratio")
plt.show()
```



```
In [20]: pca = PCA(n_components = 5)
X = pca.fit_transform(X)
pca.explained_variance_ratio_.sum()
```

Out[20]: 1.0

```
In [21]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [22]: models = pd.DataFrame(columns=["Model", "Accuracy Score"])
```

```
In [23]: log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
predictions = log_reg.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)
new_row = {"Model": "LogisticRegression", "Accuracy Score": score}
models = models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.7201492537313433

```
In [24]: rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
predictions = rfc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

new_row = {"Model": "RandomForestClassifier", "Accuracy Score": score}
models = models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.6455223880597015

```
In [25]: gbc = GradientBoostingClassifier()
gbc.fit(X_train, y_train)
predictions = gbc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

new_row = {"Model": "GradientBoostingClassifier", "Accuracy Score": score}
models = models.append(new_row, ignore_index=True)
```


Accuracy Score: 0.6977611940298507

```
In [26]: svc = SVC()
svc.fit(X_train, y_train)
predictions = svc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

new_row = {"Model": "SVC", "Accuracy Score": score}
models = models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.7313432835820896

```
In [27]: models.sort_values(by="Accuracy Score", ascending=False)
plt.figure(figsize=(12,8))
sns.barplot(x=models["Model"], y=models["Accuracy Score"])
plt.title("Models' Accuracy Scores", size=15)
plt.xticks(rotation=30)
plt.show()
```



```
In [29]: def visualize_roc_auc_curve(model, model_name):
pred_prob = model.predict_proba(X_test)
fpr, tpr, thresh = roc_curve(y_test, pred_prob[:,1], pos_label=1)

score = roc_auc_score(y_test, pred_prob[:, 1])

plt.figure(figsize=(10,8))
plt.plot(fpr, tpr, linestyle="--", color="orange", label="ROC curve (area = %0.5f)" % score)
plt.plot([0, 1], [0, 1], color="navy", lw=2, linestyle="--")

plt.title(f"{model_name} ROC Curve", size=15)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right", prop={'size': 15})
plt.show()
```


```
In [32]: tuned_models = pd.DataFrame(columns=["Model", "Accuracy Score"])
param_grid_log_reg = {"C": [0.0001, 0.001, 0.01, 0.1, 1, 10]}
grid_log_reg = GridSearchCV(LogisticRegression(), param_grid_log_reg, scoring="accuracy", cv=5, verbose=0, n_jobs=-1)

grid_log_reg.fit(X_train, y_train)
log_reg_params = grid_log_reg.best_params_
log_reg = LogisticRegression(**log_reg_params)
log_reg.fit(X_train, y_train)
predictions = log_reg.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

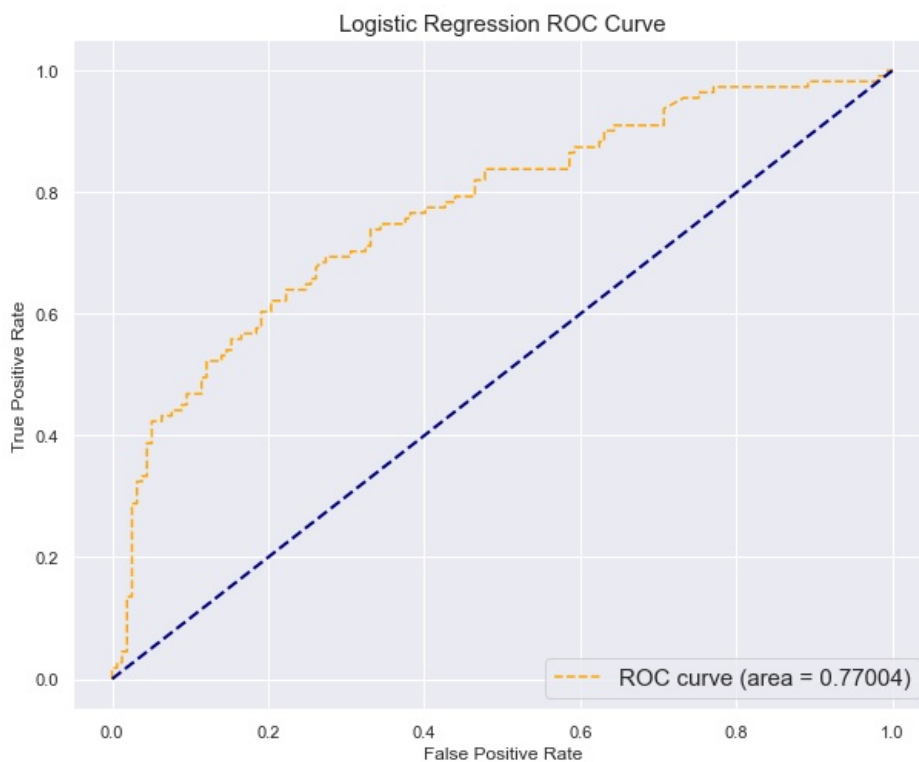
new_row = {"Model": "LogisticRegression", "Accuracy Score": score}
tuned_models = tuned_models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.7201492537313433

```
In [33]: plt.figure(figsize=(8,6))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, cmap="Blues", fmt="d")
plt.title("Confusion Matrix of Logistic Regression", size=15)
plt.show()
```



In [34]: visualize_roc_auc_curve(log_reg, "Logistic Regression")



```
In [35]: param_grid_rfc = {"min_samples_split": [2, 3, 10],
                           "min_samples_leaf": [1, 3, 10],
                           "n_estimators": [100, 200, 500]}


grid_rfc = GridSearchCV(RandomForestClassifier(), param_grid_rfc, scoring="accuracy", cv=5, verbose=0, n_jobs=-1)

grid_rfc.fit(X_train, y_train)
rfc_params = grid_rfc.best_params_
rfc = RandomForestClassifier(**rfc_params)
rfc.fit(X_train, y_train)
predictions = rfc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)
```

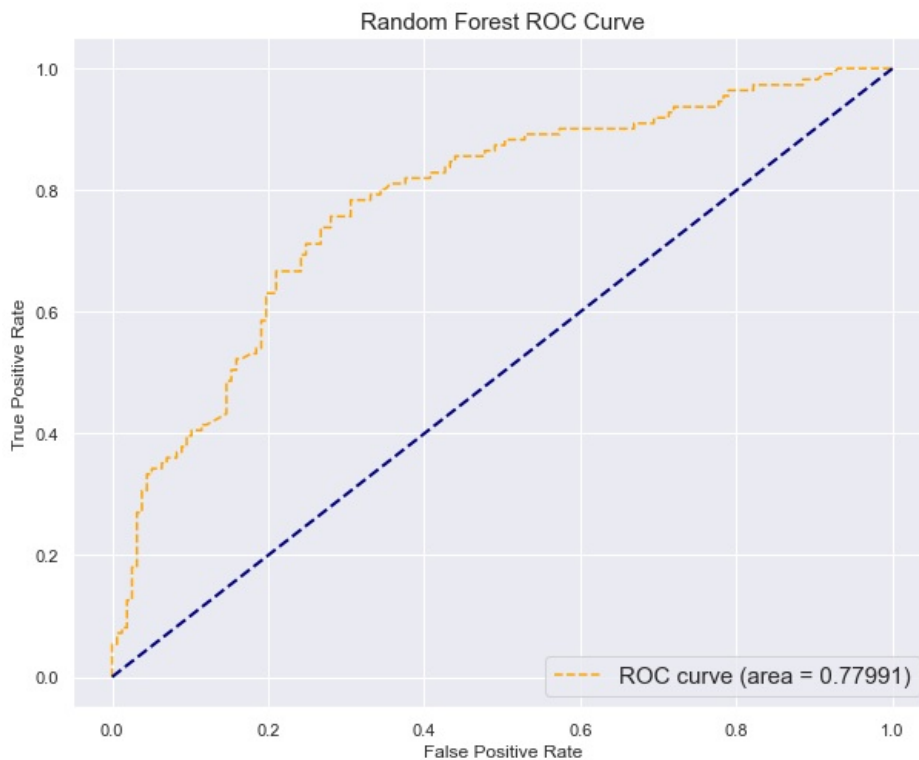
```
new_row = {"Model": "RandomForestClassifier", "Accuracy Score": score}
tuned_models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.7052238805970149

```
In [36]: plt.figure(figsize=(8,6))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, cmap="Blues", fmt='d')
plt.title("Confusion Matrix of Random Forest", size=15)
plt.show()
```



In [37]: visualize_roc_auc_curve(rfc, "Random Forest")



```
In [38]: param_grid_gbc = {'n_estimators': [100, 200, 500],
                           'learning_rate': [0.1, 0.05, 0.01],
                           'max_depth': [2, 3, 6],
                           'min_samples_leaf': [1, 2, 5]}

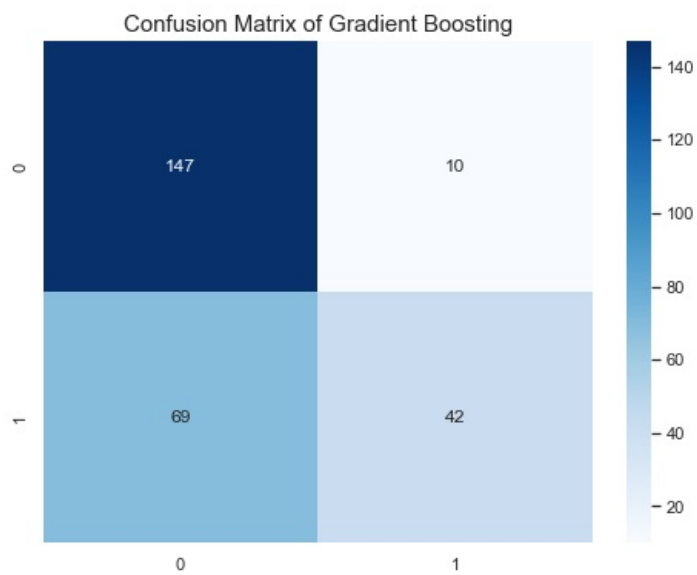
grid_gbc = GridSearchCV(GradientBoostingClassifier(), param_grid_gbc, scoring="accuracy", cv=5, verbose=0, n_jobs=-1)

grid_gbc.fit(X_train, y_train)
gbc_params = grid_gbc.best_params_
gbc = GradientBoostingClassifier(**gbc_params)
gbc.fit(X_train, y_train)
predictions = gbc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

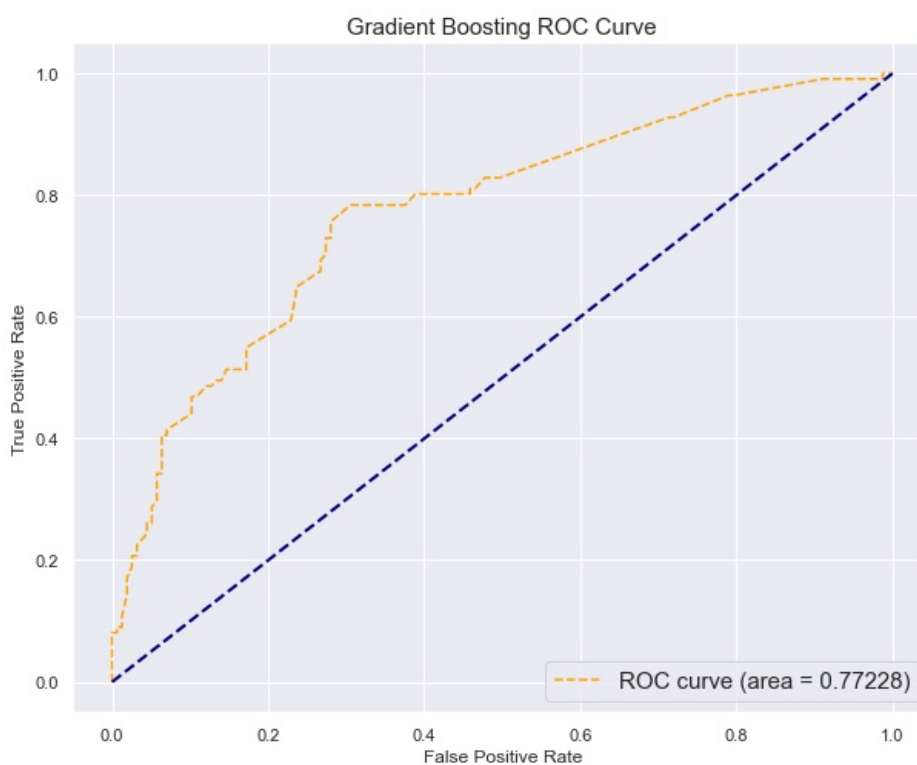
new_row = {"Model": "GradientBoostingClassifier", "Accuracy Score": score}
tuned_models = tuned_models.append(new_row, ignore_index=True)

Accuracy Score: 0.7052238805970149
```

```
In [39]: plt.figure(figsize=(8,6))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, cmap="Blues", fmt='d')
plt.title("Confusion Matrix of Gradient Boosting", size=15)
plt.show()
```



```
In [40]: visualize_roc_auc_curve(gbc, "Gradient Boosting")
```



```
In [41]: param_grid_svc = {'gamma': [ 0.001, 0.01, 0.1, 1, 10],
                          'C': [1, 10, 50, 100, 200, 300, 500, 1000]}

grid_svc = GridSearchCV(SVC(), param_grid_svc, scoring="accuracy", cv=5, verbose=0, n_jobs=-1)
grid_svc.fit(X_train, y_train)
```

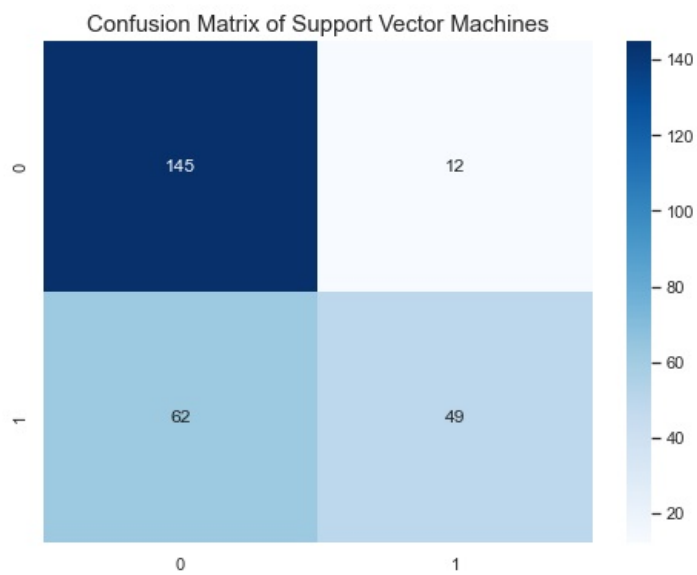
```
Out[41]: GridSearchCV(cv=5, estimator=SVC(), n_jobs=-1,
                    param_grid={'C': [1, 10, 50, 100, 200, 300, 500, 1000],
                                'gamma': [0.001, 0.01, 0.1, 1, 10]},
                    scoring='accuracy')
```

```
In [42]: svc_params = grid_svc.best_params_
svc = SVC(**svc_params)
svc.fit(X_train, y_train)
predictions = svc.predict(X_test)
score = accuracy_score(y_test, predictions)
print("Accuracy Score:", score)

new_row = {"Model": "SVC", "Accuracy Score": score}
tuned_models = tuned_models.append(new_row, ignore_index=True)
```

Accuracy Score: 0.7238805970149254

```
In [43]: plt.figure(figsize=(8,6))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True, cmap="Blues", fmt='d')
plt.title("Confusion Matrix of Support Vector Machines", size=15)
plt.show()
```

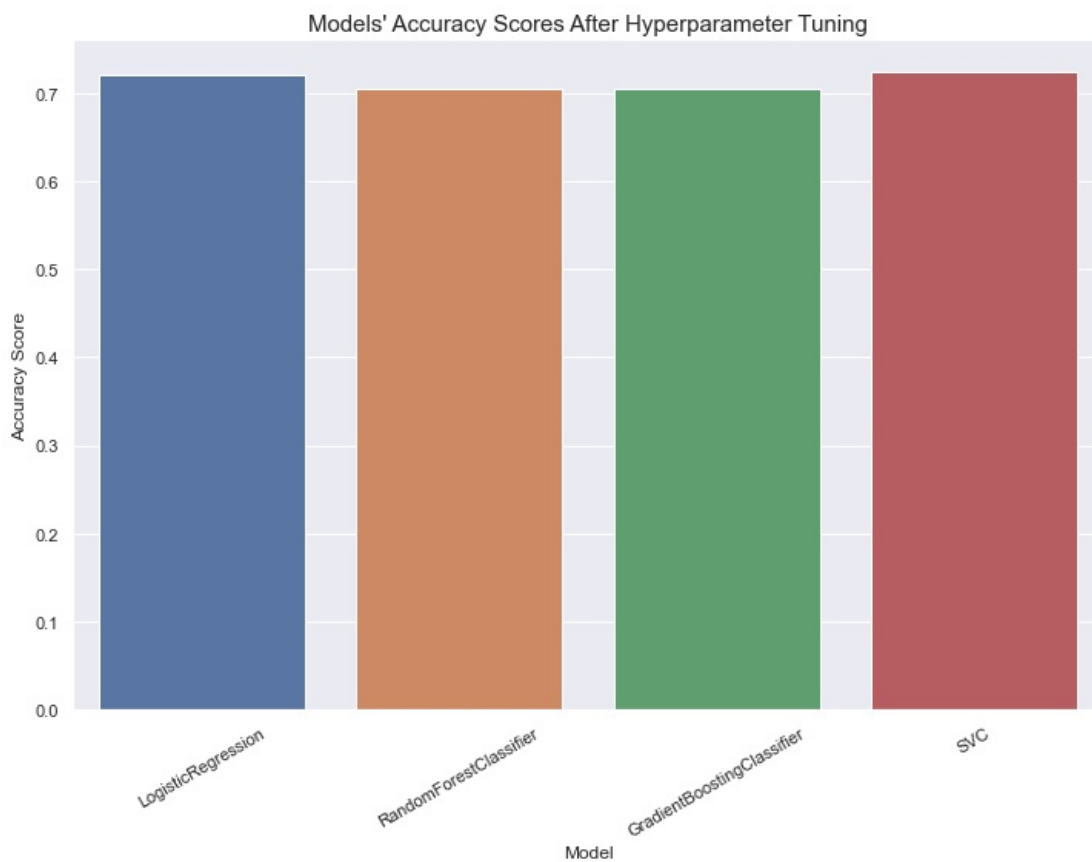



```
In [44]: tuned_models.sort_values(by="Accuracy Score", ascending=False)
```

Out[44]:

	Model	Accuracy Score
3	SVC	0.723881
0	LogisticRegression	0.720149
1	RandomForestClassifier	0.705224
2	GradientBoostingClassifier	0.705224

```
In [45]: plt.figure(figsize=(12, 8))
sns.barplot(x=tuned_models["Model"], y=tuned_models["Accuracy Score"])
plt.title("Models' Accuracy Scores After Hyperparameter Tuning", size=15)
plt.xticks(rotation=30)
plt.show()
```



```
In [ ]:
```